

---

# **CAP 5516**

## **Medical Image Computing (Spring 2025)**

**Dr. Chen Chen**  
**Associate Professor**  
**Center for Research in Computer Vision (CRCV)**  
**University of Central Florida**  
**Office: HEC 221**  
**Email: [chen.chen@crcv.ucf.edu](mailto:chen.chen@crcv.ucf.edu)**  
**Web: <https://www.crcv.ucf.edu/chenchen/>**

---



# Lecture 13

## Efficient Deep Learning (3)

# Research Challenges

---

Large models




Devices with varying  
resources




# Efficient Neural Networks Design

Credit: Vivienne Sze

## Network Pruning



## Efficient Network Architectures




[ MobileNets, ShuffleNets, AdderNet ]

## Reduce Precision

32-bit float	10100101000000000101000000000100
8-bit fixed	01100110
Binary	0

## Knowledge Distillation



Source:

<https://towardsdatascience.com/knowledge-distillation-simplified-dd4973dbc764>

# Med-DANet-V2

- Shen, Haoran, Yifu Zhang, Wenxuan Wang, Chen Chen, Jing Liu, Shanshan Song, and Jiangyun Li. "Med-DANet V2: A Flexible Dynamic Architecture for Efficient Medical Volumetric Segmentation." In Proceedings of the IEEE/CVF Winter Conference on Applications of Computer Vision, pp. 7871-7881. 2024.

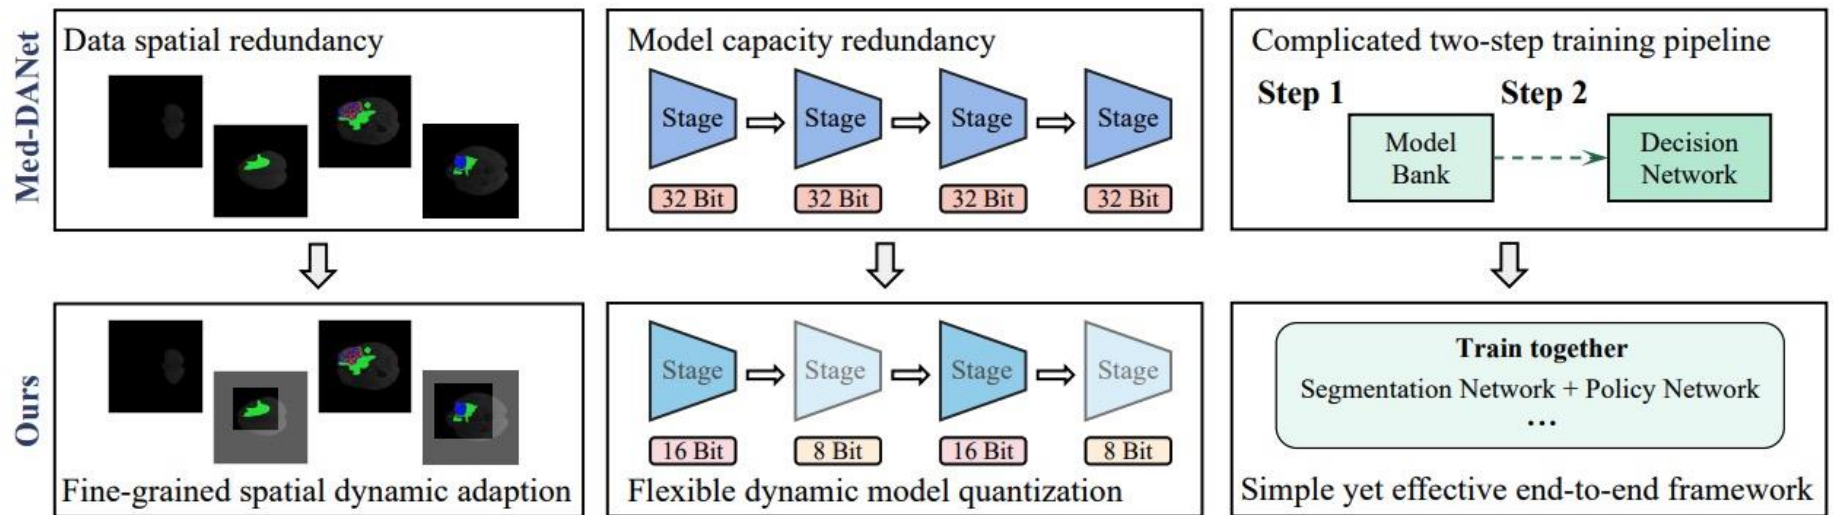
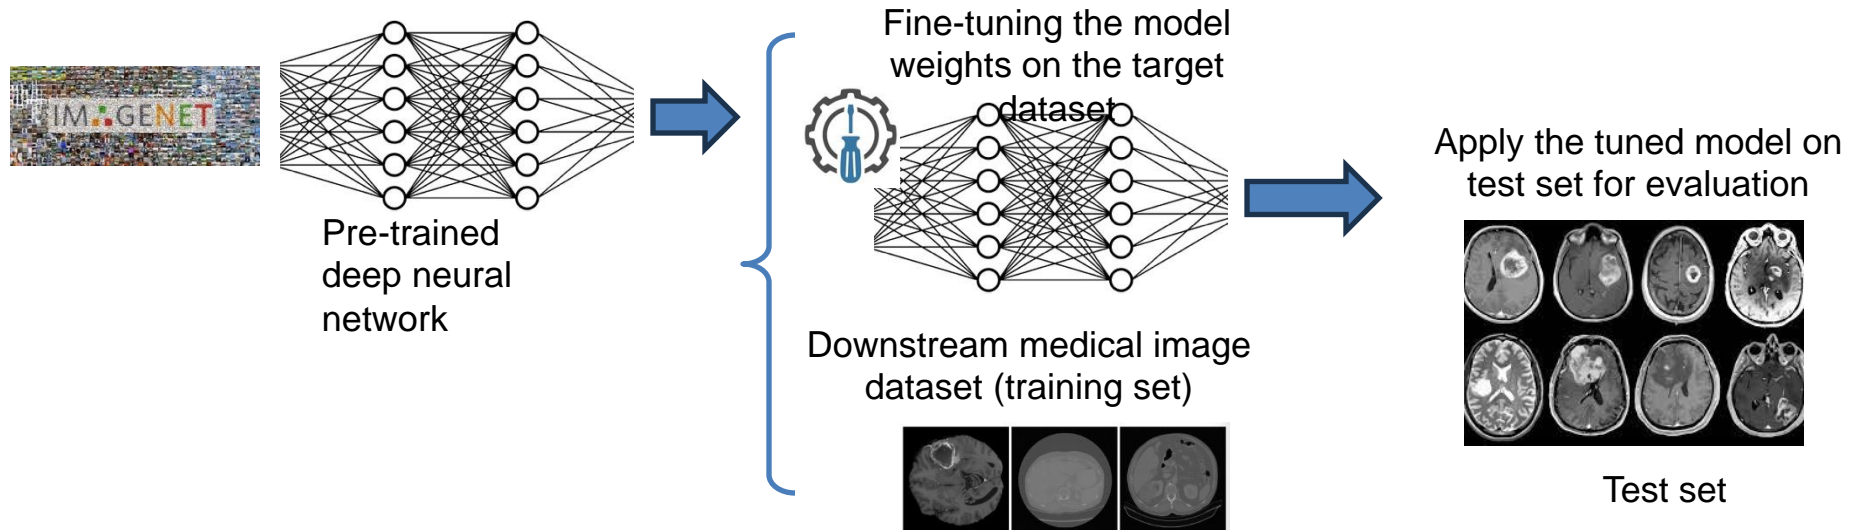


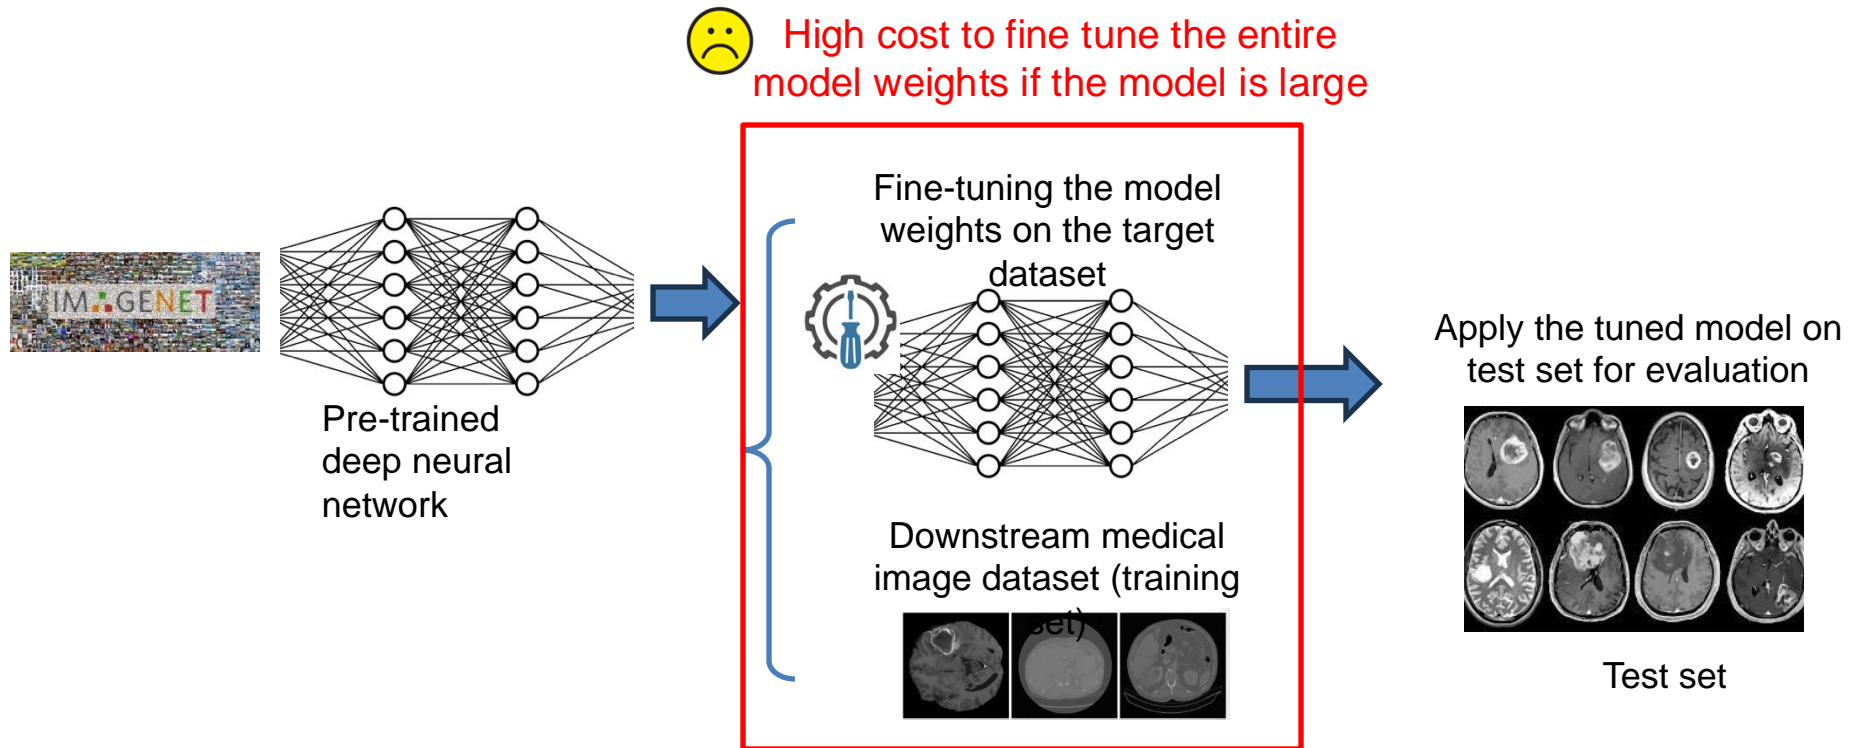
Figure 1. The comparison between the previous dynamic network Med-DANet and our proposed Med-DANet V2 (Ours).

# Motivation

- The standard “Pre-training then fine-tuning” paradigm for medical image segmentation



# Motivation



---

# Parameter Efficient Fine Tuning for Transformer models



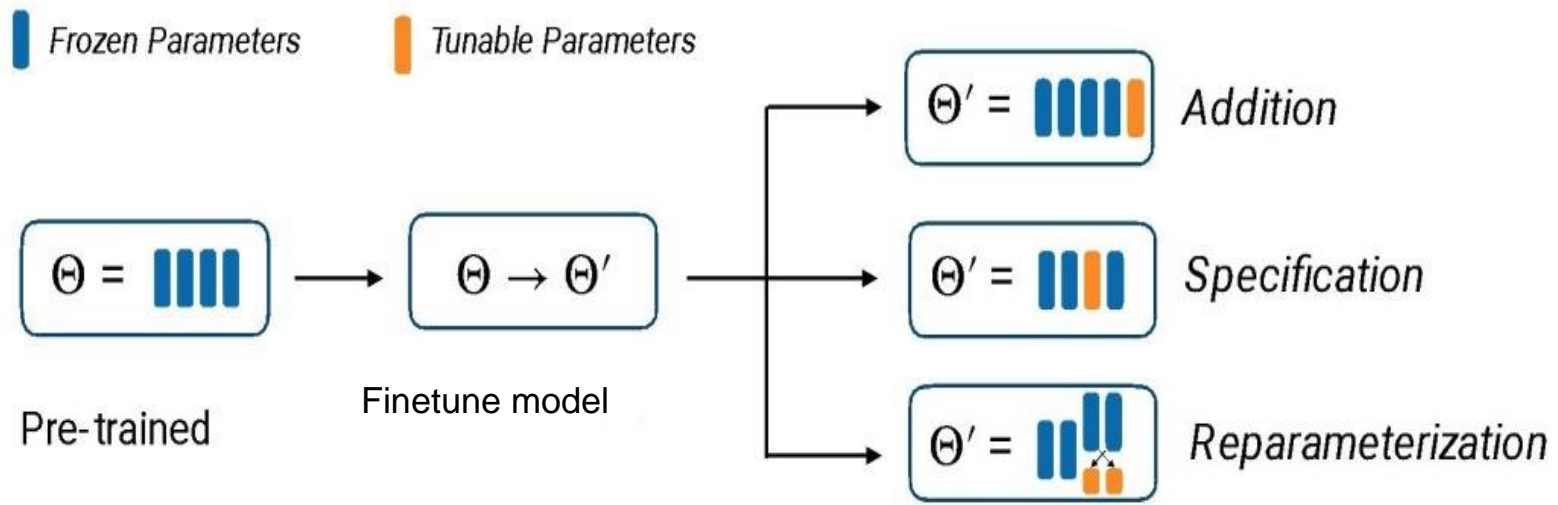
# Parameter Efficient Fine Tuning

---

- Idea: Rather than finetuning the entire model, we finetune only small amounts of weights.
- **Efficiency:** Only a small set of parameters is trained, leading to faster convergence and reduced resource usage.
- **Practical Benefits:**
  - Reduces risk of overfitting when data is scarce.
  - Enables rapid adaptation to new tasks without altering the entire network.

# Parameter Efficient Fine Tuning

- 1. Addition: What if we introduce additional trainable parameters to the neural network and just train those?
- 2. Specification: What if we pick a small subset of the parameters of the neural network and just tune those?
- 3. Reparameterization: What if we re-parameterize the model into something that is more efficient to train?



Slide Credit: Daphne Ippolito, Chenyan Xiong

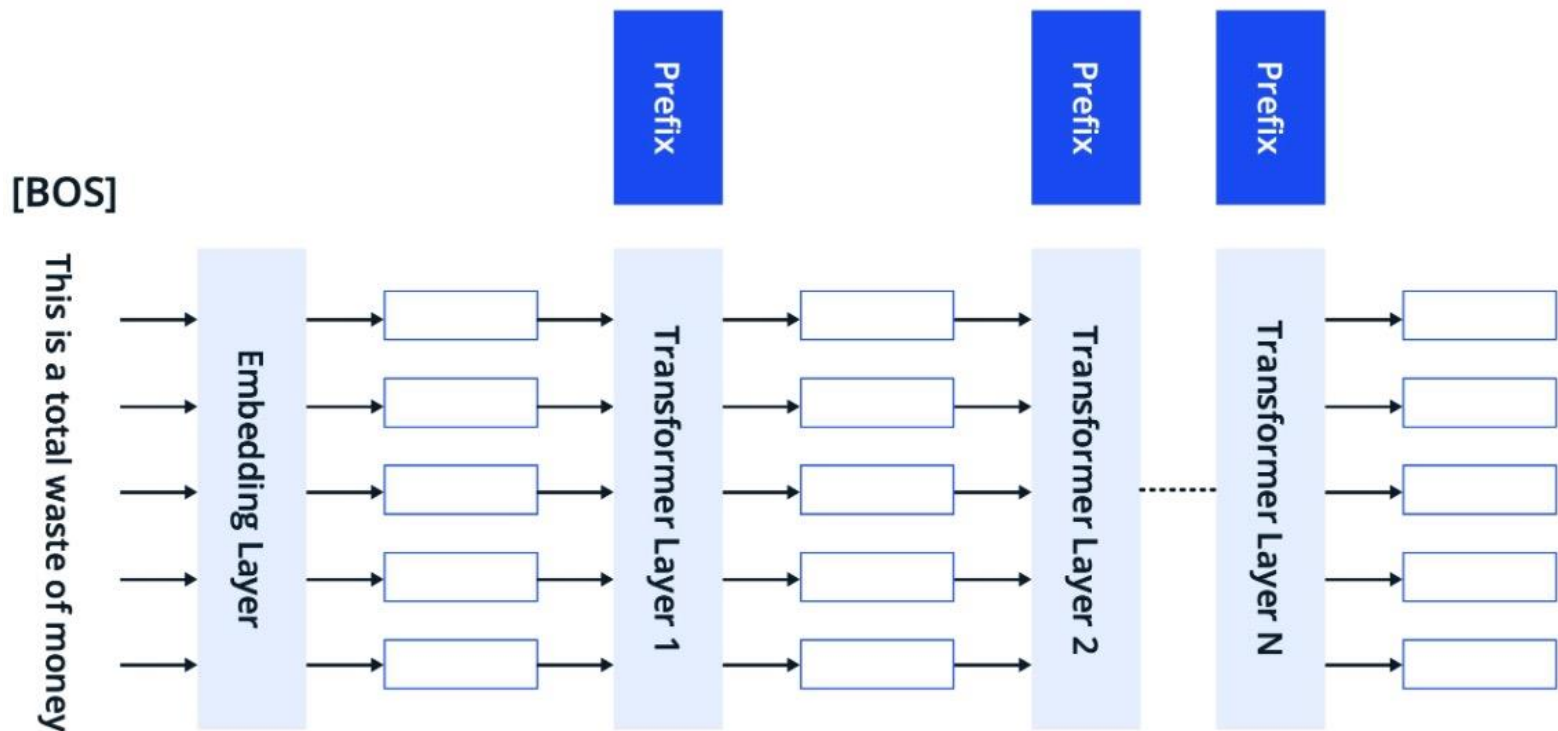
# Parameter Efficient Fine Tuning – Prefix Tuning

## Key Idea:

- Instead of updating all model weights, learn a set of additional “prefix” tokens.

## Implementation:

- **Prefix Tokens:** Learnable continuous vectors prepended to the input sequence. These prefixes are injected into every layer of the Transformer
- **Frozen Backbone:** The main pretrained model remains unchanged during fine-tuning.



# Parameter Efficient Fine Tuning – Prefix Tuning

---

## Training Dynamics

- **Frozen Model:**

The original model parameters remain frozen, ensuring that only the prefix tokens are updated.

- **Backpropagation:**

Gradients are computed only for the prefix tokens, allowing them to learn task-specific representations.

- **Outcome:**

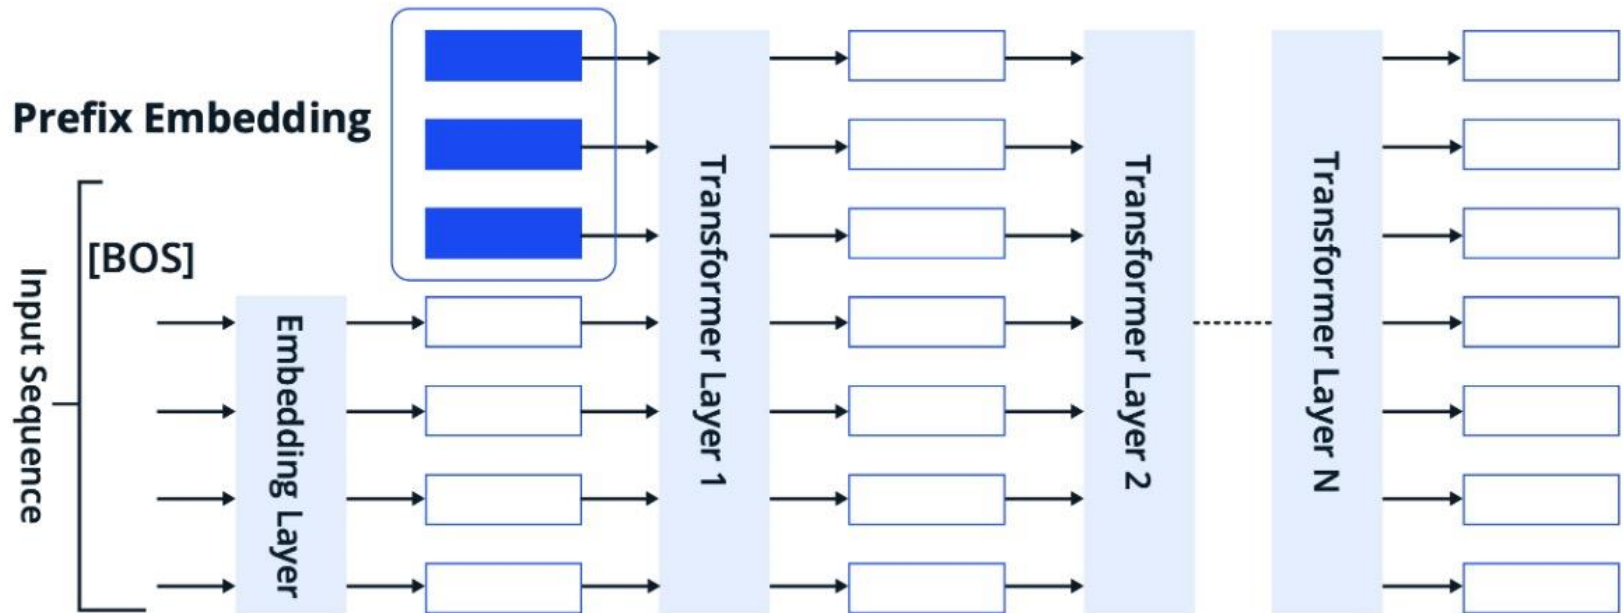
The modified attention mechanism effectively “steers” the model toward desired behaviors, achieving performance comparable to full fine-tuning but at a fraction of the cost.



# Parameter Efficient Fine Tuning – Prompt Tuning

## Implementation:

- **Prompt Tokens:** A sequence of learnable vectors, prepended to the input.
- **Frozen Backbone:** The main pretrained model remains unchanged during training.



# Parameter Efficient Fine Tuning – Prompt Tuning

---

## During Training and Inference

- Prepend these prompt tokens to every input sequence.
- The model processes the prompt tokens along with the actual input, effectively incorporating task-specific context.

## Training Dynamics

- **Frozen Model:**

Only the prompt tokens are updated; the rest of the model's parameters remain fixed.

- **Learning Process:**

The prompt tokens are optimized via backpropagation to maximize performance on the task.



# Uses Cases and Trade-offs

---

## **Prefix Tuning:**

### **Best For:**

- Tasks requiring more nuanced control over the model's internal representations.
- Scenarios where achieving near full fine-tuning performance is critical.

## **Prompt Tuning:**

### **Best For:**

- Quick adaptation to new tasks with minimal changes.
- Situations with limited computational resources and simpler integration requirements.

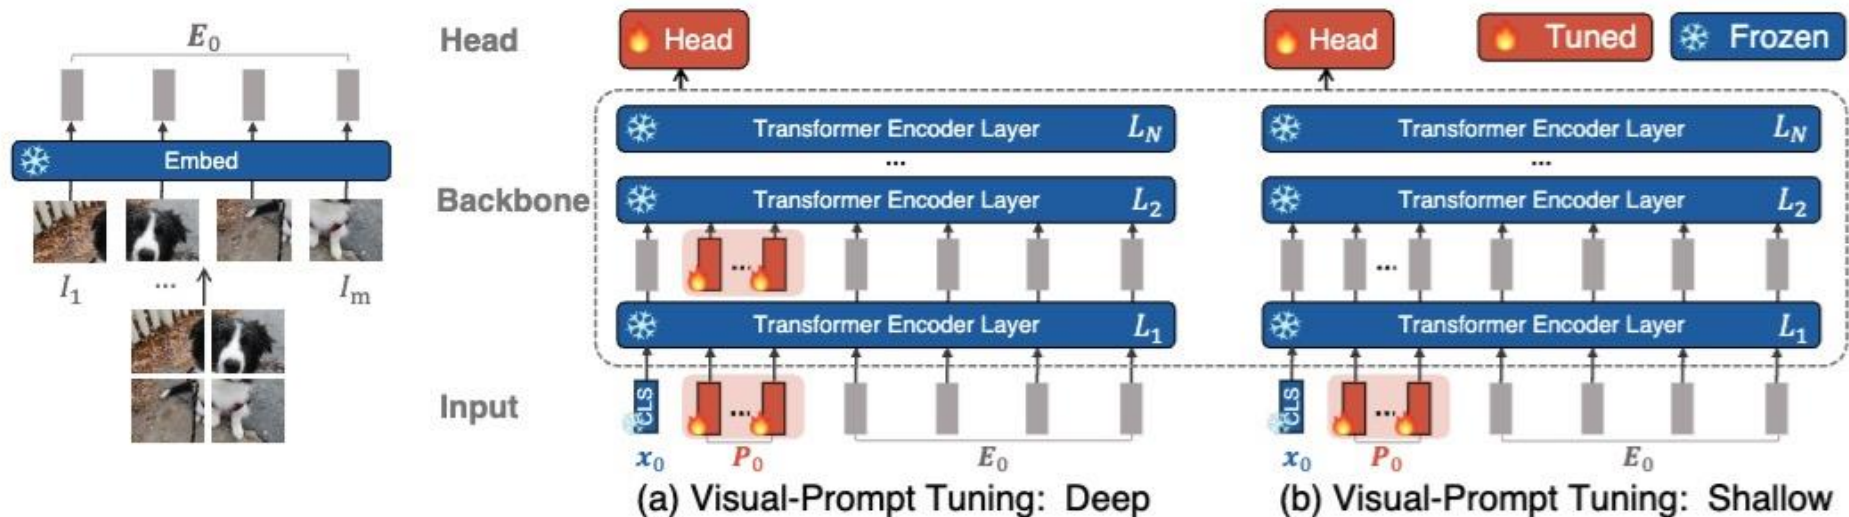
## **Trade-offs:**

### **Complexity vs. Simplicity:**

- Prefix tuning is slightly more complex but can yield finer control.
- Prompt tuning is simpler and more straightforward, but might not match the performance gains in all cases.

# Parameter Efficient Fine Tuning – Visual Prompt Tuning

- Visual Prompt Tuning



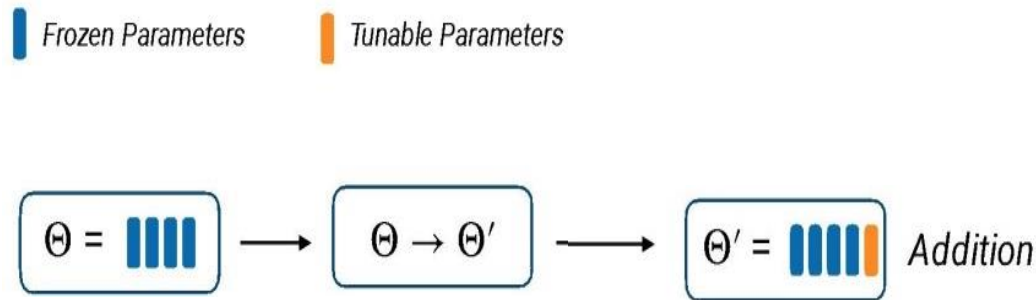
**Fig. 2.** Overview of our proposed Visual-Prompt Tuning. We explore two variants: (a) prepend a set of learnable parameters to each Transformer encoder layer's input (VPT-DEEP); (b) only insert the prompt parameters to the first layer's input (VPT-SHALLOW). During training on downstream tasks, only the parameters of prompts and linear head are updated while the whole Transformer encoder is frozen.

Jia, M., Tang, L., Chen, B. C., Cardie, C., Belongie, S., Hariharan, B., & Lim, S. N. (2022, October). Visual prompt tuning. In European Conference on Computer Vision (pp. 709-727). Cham: Springer Nature Switzerland.



# Parameter Efficient Fine Tuning - Adaptor

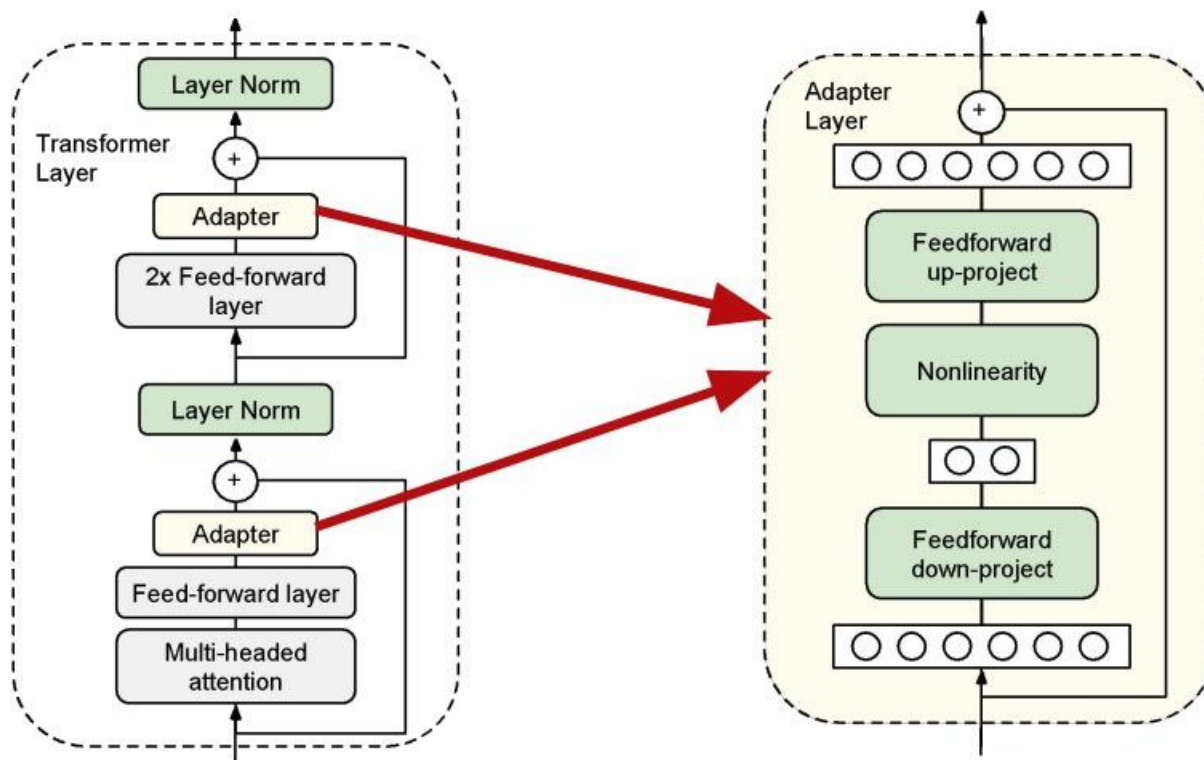
- Addition: What if we introduce additional (a small amount) trainable parameters to the neural network and just train those?



- Adapters are new modules are added between layers of a pre-trained network.
- The original model weights are fixed; just the adapter modules are tuned.

Houlsby, Neil, et al. "Parameter-efficient transfer learning for NLP." International Conference on Machine Learning. PMLR, 2019.

# Parameter Efficient Fine Tuning - Adaptor



The adapter consists of a bottleneck which contains few parameters relative to the attention and feedforward layers in the original model.

Only the adaptor layers (i.e., a small set of parameters) are trained on the downstream data

Houlsby, Neil, et al. "Parameter-efficient transfer learning for NLP." International Conference on Machine Learning. PMLR, 2019.



# Parameter Efficient Fine Tuning - LoRA

## LoRA: LOW-RANK ADAPTATION OF LARGE LANGUAGE MODELS

Edward Hu\*   Yelong Shen\*   Phillip Wallis   Zeyuan Allen-Zhu  
Yuanzhi Li   Shean Wang   Lu Wang   Weizhu Chen

Microsoft Corporation

{edwardhu, yeshe, phwallis, zeyuana,  
yuanzhil, swang, luw, wzchen}@microsoft.com

yuanzhil@andrew.cmu.edu

(Version 2)

### ABSTRACT

An important paradigm of natural language processing consists of large-scale pre-training on general domain data and adaptation to particular tasks or domains. As we pre-train larger models, full fine-tuning, which retrains all model parameters, becomes less feasible. Using GPT-3 175B as an example – deploying independent instances of fine-tuned models, each with 175B parameters, is prohibitively expensive. We propose **Low-Rank Adaptation**, or LoRA, which freezes the pre-trained model weights and injects trainable rank decomposition matrices into each layer of the Transformer architecture, greatly reducing the number of trainable parameters for downstream tasks. Compared to GPT-3 175B fine-tuned with Adam, LoRA can reduce the number of trainable parameters by 10,000 times and the GPU memory requirement by 3 times. LoRA performs on-par or better than fine-tuning in model quality on RoBERTa, DeBERTa, GPT-2, and GPT-3, despite having fewer trainable parameters, a higher training throughput, and, unlike adapters, *no additional inference latency*. We also provide an empirical investigation into rank-deficiency in language model adaptation, which sheds light on the efficacy of LoRA. We release a package that facilitates the integration of LoRA with PyTorch models and provide our implementations and model checkpoints for RoBERTa, DeBERTa, and GPT-2 at <https://github.com/microsoft/LoRA>.

Hu, Edward J., et al. "Lora: Low-rank adaptation of large language models." *arXiv preprint arXiv:2106.09685* (2021).



# Parameter Efficient Fine Tuning - LoRA

## LoRA: LOW-RANK ADAPTATION OF LARGE LANGUAGE MODELS

Edward Hu\*   Yelong Shen\*   Phillip Wallis   Zeyuan Allen-Zhu  
Yuanzhi Li   Shean Wang   Lu Wang   Weizhu Chen

Microsoft Corporation

{edwardhu, yeshe, phwallis, zeyuana,  
yuanzhil, swang, luw, wzchen}@microsoft.com

yuanzhil@andrew.cmu.edu

(Version 2)

### ABSTRACT

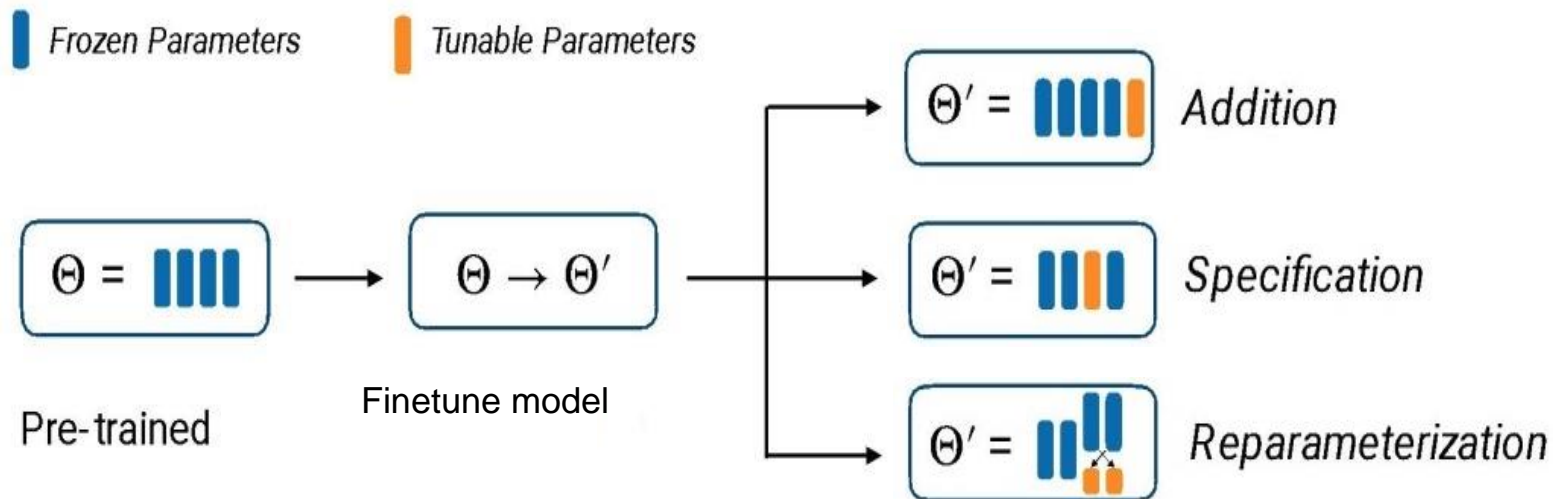
An important paradigm of natural language processing consists of large-scale pre-training on general domain data and adaptation to particular tasks or domains. As we pre-train larger models, full fine-tuning, which retrain all model parameters, becomes less feasible. Using GPT-3 175B as an example – deploying independent instances of fine-tuned models, each with 175B parameters, is prohibitively expensive. We propose **Low-Rank Adaptation**, or LoRA, which freezes the pre-trained model weights and injects trainable rank decomposition matrices into each layer of the Transformer architecture, greatly reducing the number of trainable parameters for downstream tasks. Compared to GPT-3 175B fine-tuned with Adam, LoRA can reduce the number of trainable parameters by 10,000 times and the GPU memory requirement by 3 times. LoRA performs on-par or better than fine-tuning in model quality on RoBERTa, DeBERTa, GPT-2, and GPT-3, despite having fewer trainable parameters, a higher training throughput, and, unlike adapters, *no additional inference latency*. We also provide an empirical investigation into rank-deficiency in language model adaptation, which sheds light on the efficacy of LoRA. We release a package that facilitates the integration of LoRA with PyTorch models and provide our implementations and model checkpoints for RoBERTa, DeBERTa, and GPT-2 at <https://github.com/microsoft/LoRA>.

Hu, Edward J., et al. "Lora: Low-rank adaptation of large language models." *arXiv preprint arXiv:2106.09685* (2021).



# Recall: Parameter Efficient Fine Tuning

- 1. Addition: What if we introduce additional trainable parameters to the neural network and just train those?
- 2. Specification: What if we pick a small subset of the parameters of the neural network and just tune those?
- 3. Reparameterization: What if we re-parameterize the model into something that is more efficient to train?



# Parameter Efficient Fine Tuning - LoRA

---

- Previous study shows that
  - Pre-trained LLMs have a “low intrinsic dimension”
  - LLMs can still learn efficiently despite a low-dim reparametrization



# Parameter Efficient Fine Tuning - LoRA

## Key Idea

- Keep the original pretrained parameters  $\mathbf{W}_0$  fixed during fine-tuning
- Learn an additive modification to those parameters  $\Delta\mathbf{W}$
- Define  $\Delta\mathbf{W}$  via a low rank decomposition:


$$\Delta\mathbf{W} = \mathbf{BA}$$

where  $\mathbf{BA}$  has rank  $r$ , which is **much less** than the input dimension  $k$  or the output dimension  $d$

Standard Linear Layer

$$\mathbf{z} = \mathbf{W}_0\mathbf{x}$$


$$\mathbf{W}_0 \in \mathbb{R}^{d \times k}, \mathbf{x} \in \mathbb{R}^k, \mathbf{z} \in \mathbb{R}^d$$



LoRA Linear Layer

$$\begin{aligned}\mathbf{z} &= \mathbf{W}_0\mathbf{x} + \mathbf{BAx} \\ &= (\mathbf{W}_0 + \mathbf{BA})\mathbf{x}\end{aligned}$$

$$\begin{aligned}\mathbf{W}_0 &\in \mathbb{R}^{d \times k}, \\ \mathbf{A} &\in \mathbb{R}^{r \times k}, \mathbf{B} \in \mathbb{R}^{d \times r} \\ \text{where } r &\ll \min(d, k)\end{aligned}$$



Slide credit: Matt Gormley

# Parameter Efficient Fine Tuning - LoRA

---

- **Low-Rank Decomposition:**

- When you decompose a weight update  $\Delta W$  into two matrices  $A \in \mathbb{R}^{d \times r}$  and  $B \in \mathbb{R}^{r \times k}$ , the product  $A \times B$  is a matrix that has at most rank  $r$ .
- In linear algebra, the rank of a matrix is the number of its linearly independent rows or columns. By constraining the update to be of rank  $r$ , you are effectively reducing the degrees of freedom in the update.

- **Control over Complexity:**

- A lower  $r$  means that the update is captured in a smaller subspace, which forces the model to make only the most essential adjustments needed for the task.
- This is key for efficiency, as it drastically reduces the number of trainable parameters while still allowing meaningful adaptation.



# Parameter Efficient Fine Tuning - LoRA

## Initialization

- We initialize the trainable parameters:  
 $A_{ij} \sim \mathcal{N}(0, \sigma^2), \forall i, j$   
 $B = 0$
- This ensures that, at the start of fine tuning, the parameters have their pretrained values:


$$\Delta W = BA = 0$$

$$W_0 + BA = W_0$$

Standard Linear Layer

$$z = W_0 x$$

$$W_0 \in \mathbb{R}^{d \times k}, x \in \mathbb{R}^k, z \in \mathbb{R}^d$$




LoRA Linear Layer

$$z = W_0 x + BAx$$
$$= (W_0 + BA)x$$

$$W_0 \in \mathbb{R}^{d \times k},$$
$$A \in \mathbb{R}^{r \times k}, B \in \mathbb{R}^{d \times r}$$

where  $r \ll \min(d, k)$



Slide credit: Matt Gormley



# Parameter Efficient Fine Tuning - LoRA

- To get a Standard Linear Layer with parameters  $\mathbf{W}$  that includes our LoRA fine tuning:

$$\mathbf{W} \leftarrow \mathbf{W}_0 + \mathbf{BA}$$

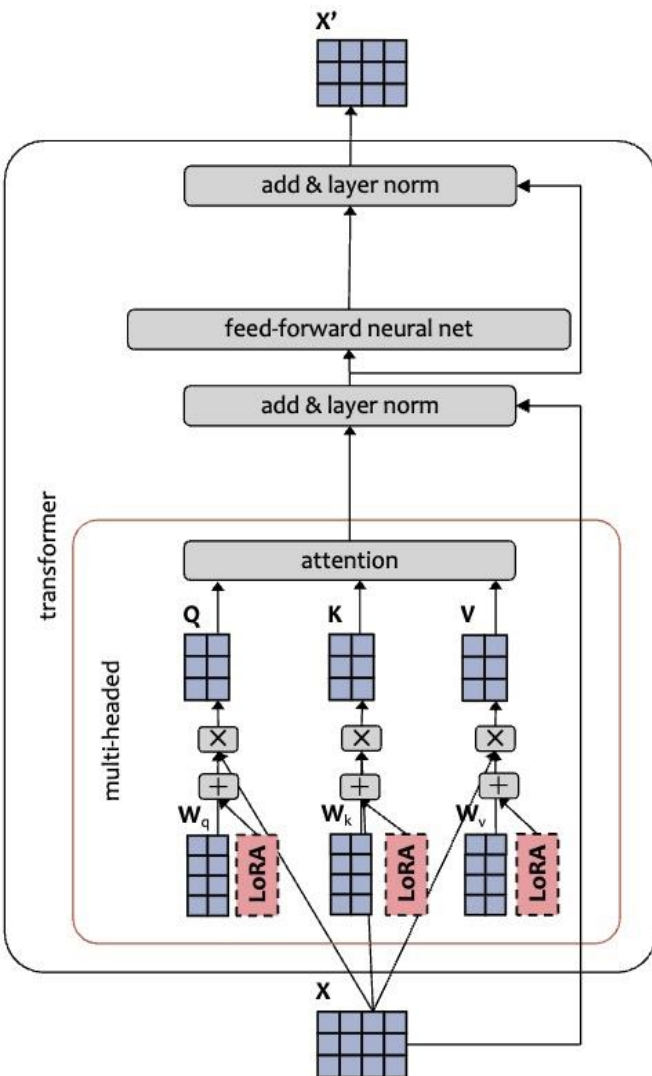
- To remove the LoRA fine tuning from that Standard Linear Layer:

$$\mathbf{W} \leftarrow \mathbf{W} - \mathbf{BA} = \mathbf{W}_0$$



Slide credit: Matt Gormley

# LoRA for Transformer




- LoRA linear layers could replace every linear layer in the Transformer layer
- But the original paper only applies LoRA to the attention weights

$$\begin{aligned} z &= W_0 x + B A x \\ &= (W_0 + B A) x \end{aligned}$$

$$\begin{aligned} W_0 &\in \mathbb{R}^{d \times k}, \\ A &\in \mathbb{R}^{r \times k}, B \in \mathbb{R}^{d \times r} \\ \text{where } r &\ll \min(d, k) \end{aligned}$$

LoRA Linear Layer



# LoRA for Transformer

- LoRA linear layers could replace every linear layer in the Transformer layer
- But the original paper only applies LoRA to the attention weights
- Empirically, for GPT-3, they also find that it is most efficient to include LoRA **only on the Q and V** linear layers:

	# of Trainable Parameters = 18M						
Weight Type Rank $r$	$W_q$ 8	$W_k$ 8	$W_v$ 8	$W_o$ 8	$W_q, W_k$ 4	$W_q, W_v$ 4	$W_q, W_k, W_v, W_o$ 2
WikiSQL ( $\pm 0.5\%$ )	70.4	70.0	73.0	73.2	71.4	<b>73.7</b>	<b>73.7</b>
MultiNLI ( $\pm 0.1\%$ )	91.0	90.8	91.0	91.3	91.3	91.3	<b>91.7</b>

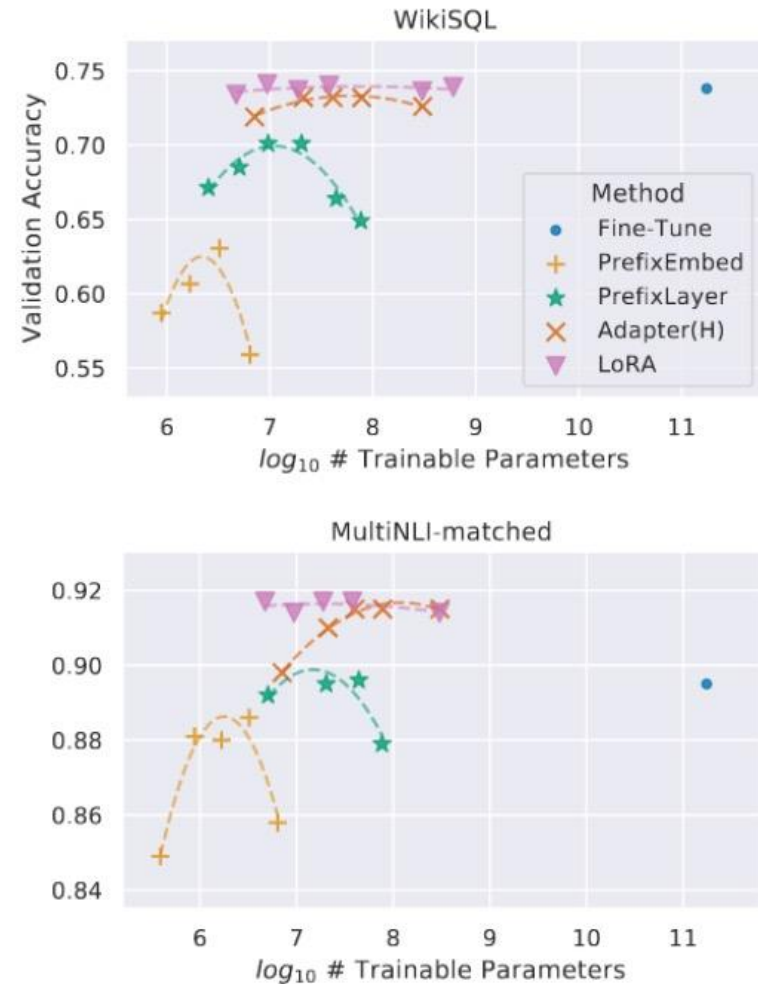
Table 5: Validation accuracy on WikiSQL and MultiNLI after applying LoRA to different types of attention weights in GPT-3, given the same number of trainable parameters. Adapting both  $W_q$  and  $W_v$  gives the best performance overall. We find the standard deviation across random seeds to be consistent for a given dataset, which we report in the first column.

Slide credit: Matt Gormley

# LoRA Results

## Takeaways

- Applied to GPT-3, LoRA achieves performance almost as good as full fine-tuning, but with far fewer parameters
- On some tasks it even outperforms full fine-tuning
- For some datasets a rank of  $r=1$  is sufficient



Slide credit: Matt Gormley

# LoRA Variants

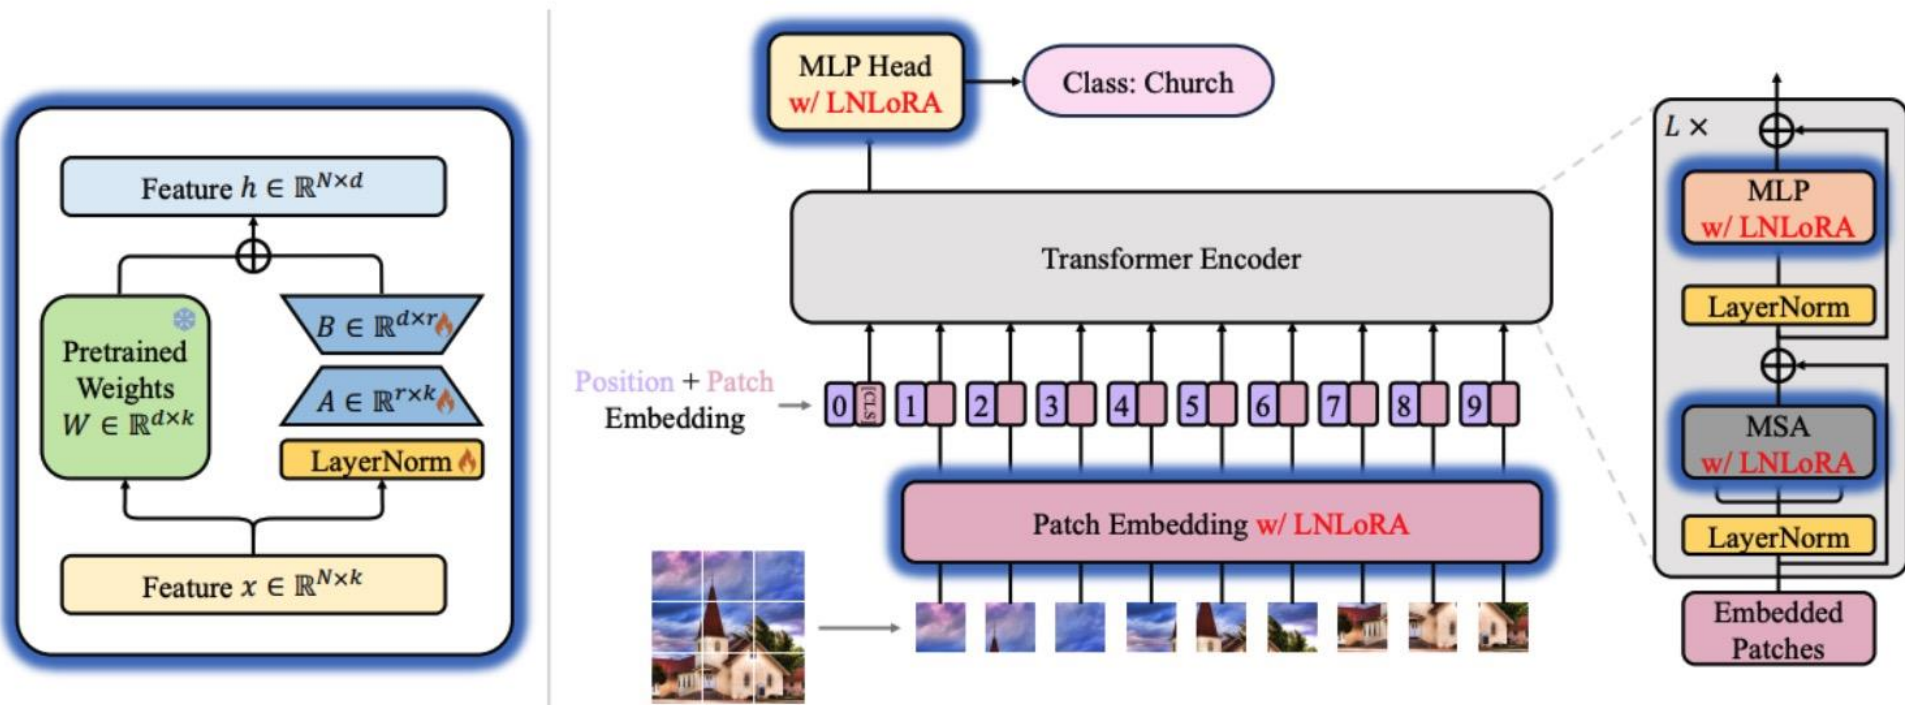
---

- LongLoRA [Chen et al., 2024]
  - Sparse Local attention to support longer context length during finetuning
- LoRA+ [Hayou et al., 2024]
  - different learning rates for the LoRA adapter matrices A and B improves finetuning speed
- DyLoRA [Valipou et al., 2023]
  - selects rank without requiring multiple runs of training




# PEFT for Vision Transformer

- Since Vision Transformer is just another transformer model, we can apply LoRA directly to it
- (LNLoRA is just a variant that includes LayerNorm in the LoRALinear module.)



Yuan, Zheng, Jie Zhang, and Shiguang Shan. "Fullora-at: Efficiently boosting the robustness of pretrained vision transformers." *arXiv preprint arXiv:2401.01752* (2024).

# PEFT for Vision Transformer



- For various computer vision tasks, parameter efficient transfer-learning (PETL) is sometimes **better** than full fine-tuning!
- VTAB-1k is a collection of 19 different vision tasks; here we're seeing average performance across tasks
- (FacT is another low-rank method capable of dramatically reducing the number of parameters tuned.)

Jie, Shibo, and Zhi-Hong Deng. "Fact: Factor-tuning for lightweight adaptation on vision transformer." *Proceedings of the AAAI conference on artificial intelligence*. Vol. 37. No. 1. 2023.

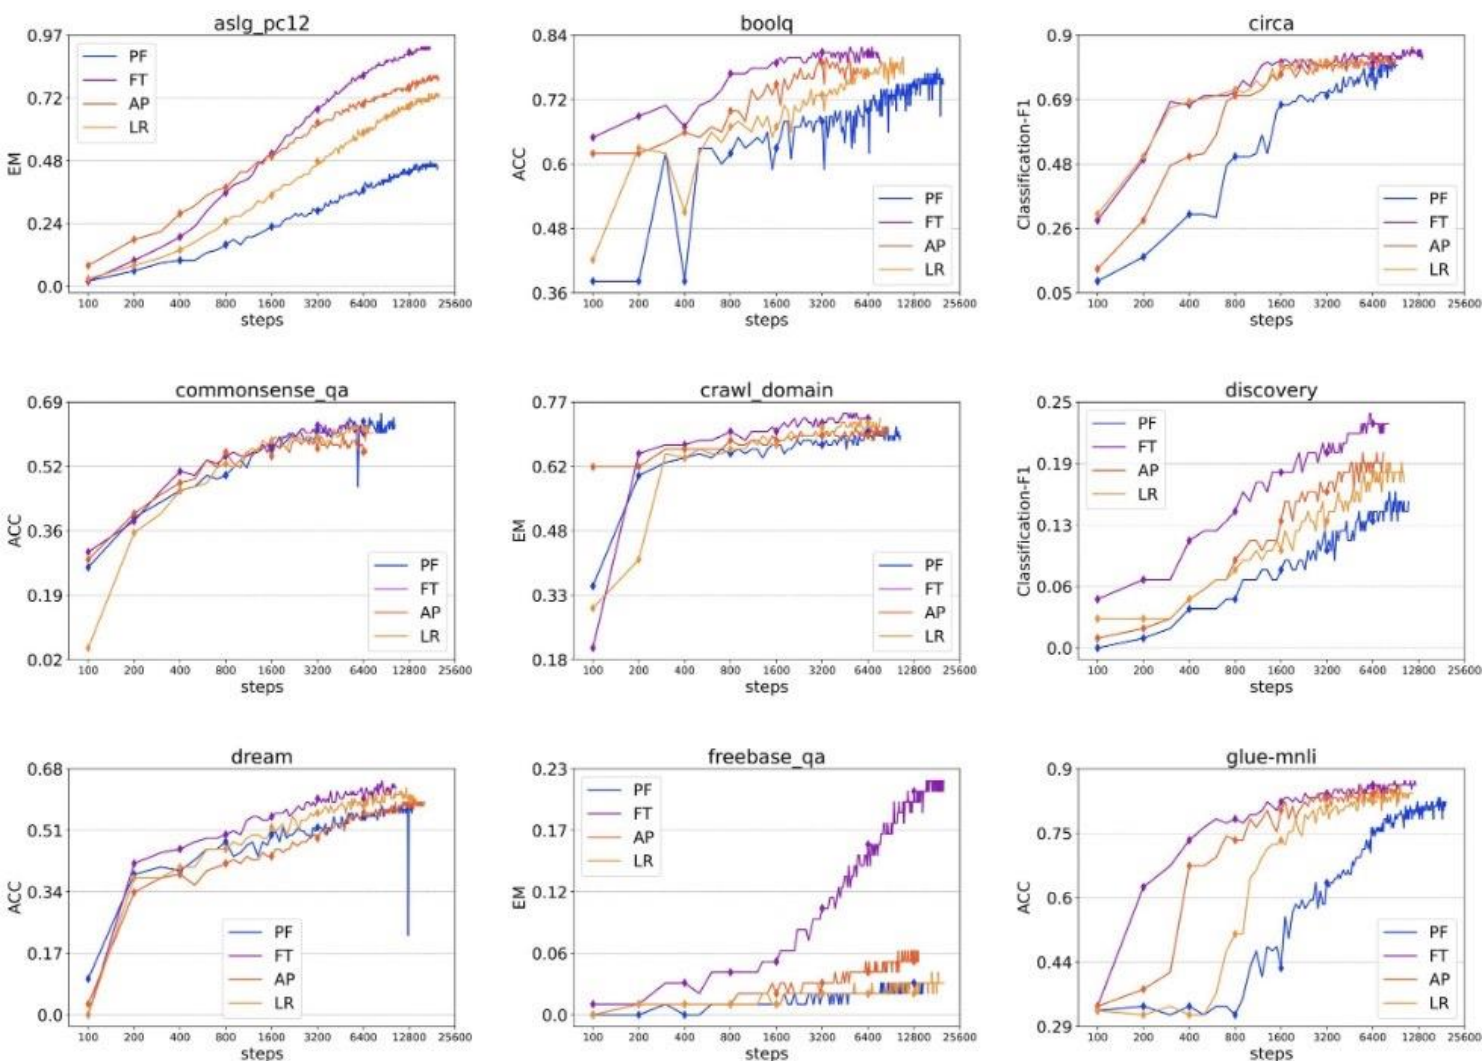


# Summary

green color : tunable parameters and modules

Name & Refs	Method	#Params
SEQUENTIAL ADAPTER Houlsby et al. (2019)	$\text{LayerNorm}(\mathbf{X} + \mathbf{H}(\mathbf{X})) \rightarrow \text{LayerNorm}(\mathbf{X} + \text{ADT}(\mathbf{H}(\mathbf{X})))$	$L \times 2 \times (2d_h d_m)$
COMPACTER Mahabadi et al. (2021a)	$\text{LayerNorm}(\mathbf{X} + \mathbf{F}(\mathbf{X})) \rightarrow \text{LayerNorm}(\mathbf{X} + \text{ADT}(\mathbf{F}(\mathbf{X})))$	$L \times 2 \times (2(d_h + d_m))$
ADAPTERDROP Rücklé et al. (2021)	$\text{ADT}(\mathbf{X}) = \mathbf{X} + \sigma(\mathbf{X}\mathbf{W}_{d_h \times d_m})\mathbf{W}_{d_m \times d_h}, \sigma = \text{activation}$	$(L - n) \times 2 \times (2d_h d_m)$
PARALLEL ADAPTER He et al. (2022)	$\text{LayerNorm}(\mathbf{X} + \mathbf{H}(\mathbf{X})) \rightarrow \text{LayerNorm}(\mathbf{X} + \text{ADT}(\mathbf{X}) + \mathbf{H}(\mathbf{X}))$ $\text{LayerNorm}(\mathbf{X} + \mathbf{F}(\mathbf{X})) \rightarrow \text{LayerNorm}(\mathbf{X} + \text{ADT}(\mathbf{X}) + \mathbf{F}(\mathbf{X}))$ $\text{ADT}(\mathbf{X}) = \sigma(\mathbf{X}\mathbf{W}_{d_h \times d_m})\mathbf{W}_{d_m \times d_h}, \sigma = \text{activation}$	$L \times 2 \times (2d_h d_m)$
ADAPTERBIAS	$\text{LayerNorm}(\mathbf{X} + \mathbf{F}(\mathbf{X})) \rightarrow \text{LayerNorm}(\text{ADT}(\mathbf{X}) + \mathbf{F}(\mathbf{X}))$ $\text{ADT}(\mathbf{X}) = \mathbf{X}\mathbf{W}_{d_h \times 1}\mathbf{W}_{1 \times d_h}$	$L \times 2 \times d_h$
PREFIX-TUNING Li & Liang (2021)	$\mathbf{H}_i = \text{ATT}(\mathbf{X}\mathbf{W}_q^{(i)}, [\text{MLP}_k^{(i)}(\mathbf{P}'_k) : \mathbf{X}\mathbf{W}_k^{(i)}], [\text{MLP}_v^{(i)}(\mathbf{P}'_v) : \mathbf{X}\mathbf{W}_v^{(i)}])$ $\text{MLP}^{(i)}(\mathbf{X}) = \sigma(\mathbf{X}\mathbf{W}_{d_m \times d_m})\mathbf{W}_{d_m \times d_h}^{(i)}$ $\mathbf{P}' = \mathbf{W}_{n \times d_m}$	$n \times d_m + d_m^2 + L \times 2 \times d_h d_m$
LoRA Hu et al. (2021a)	$\mathbf{H}_i = \text{ATT}(\mathbf{X}\mathbf{W}_q^{(i)}, \text{ADT}_k(\mathbf{X}) + \mathbf{X}\mathbf{W}_k^{(i)}, \text{ADT}_v(\mathbf{X}) + \mathbf{X}\mathbf{W}_v^{(i)})$ $\text{ADT}(\mathbf{X}) = \mathbf{X}\mathbf{W}_{d_h \times d_m}\mathbf{W}_{d_m \times d_h}$	$L \times 2 \times (2d_h d_m)$
BITFIT Zaken et al. (2021)	$f(\mathbf{X}) \rightarrow f(\mathbf{X}) + \mathbf{B}, \text{ for all function } f$	$L \times (7 \times d_h + d_m)$

# Results



PF: prefix tuning  
FT: full fine-tuning  
AP: adapter  
LR: LoRA

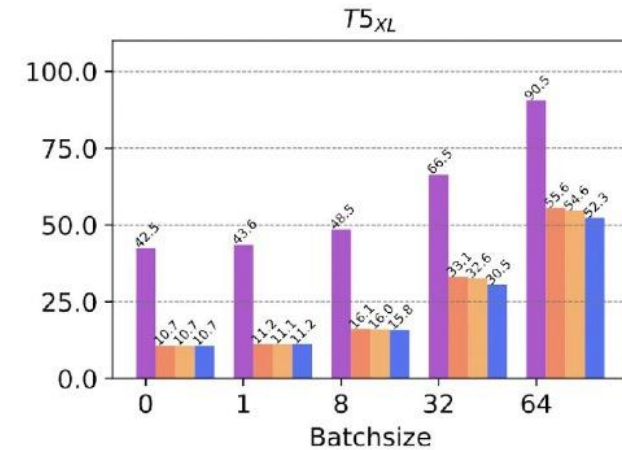


**This survey overall found:**  
Full fine-tuning >  
LoRA >  
Adapters >  
Prefix Tuning >  
Prompt Tuning

**In terms of performance.**

*Plots for many more tasks  
can be found in the paper.*



# Memory Usage



FT: full fine-tuning


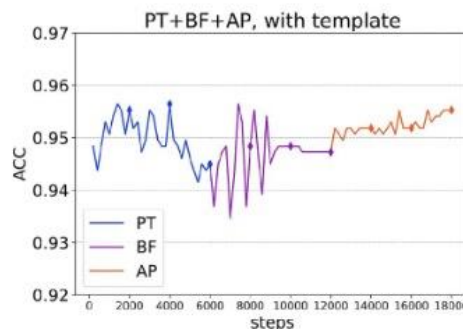




AP: adapter

LR: LoRA

BF: BitFit

Prompt tuning and prefix tuning not included because they use the same amount of memory as full fine-tuning

# Combine Different Methods



FT: BitFit

AP: adapter

PT: prompt tuning

Results on SST-2  
sentiment  
classification

# Surveys

---

- Han, Zeyu, Chao Gao, Jinyang Liu, Jeff Zhang, and Sai Qian Zhang. "Parameter-efficient fine-tuning for large models: A comprehensive survey." *arXiv preprint arXiv:2403.14608* (2024).
- Xu, Lingling, et al. "Parameter-efficient fine-tuning methods for pretrained language models: A critical review and assessment." *arXiv preprint arXiv:2312.12148* (2023).
- Xin, Yi, et al. "Parameter-efficient fine-tuning for pre-trained vision models: A survey." *arXiv preprint arXiv:2402.02242* (2024).
- Ding, Ning, et al. "Delta tuning: A comprehensive study of parameter efficient methods for pre-trained language models." *arXiv preprint arXiv:2203.06904* (2022).



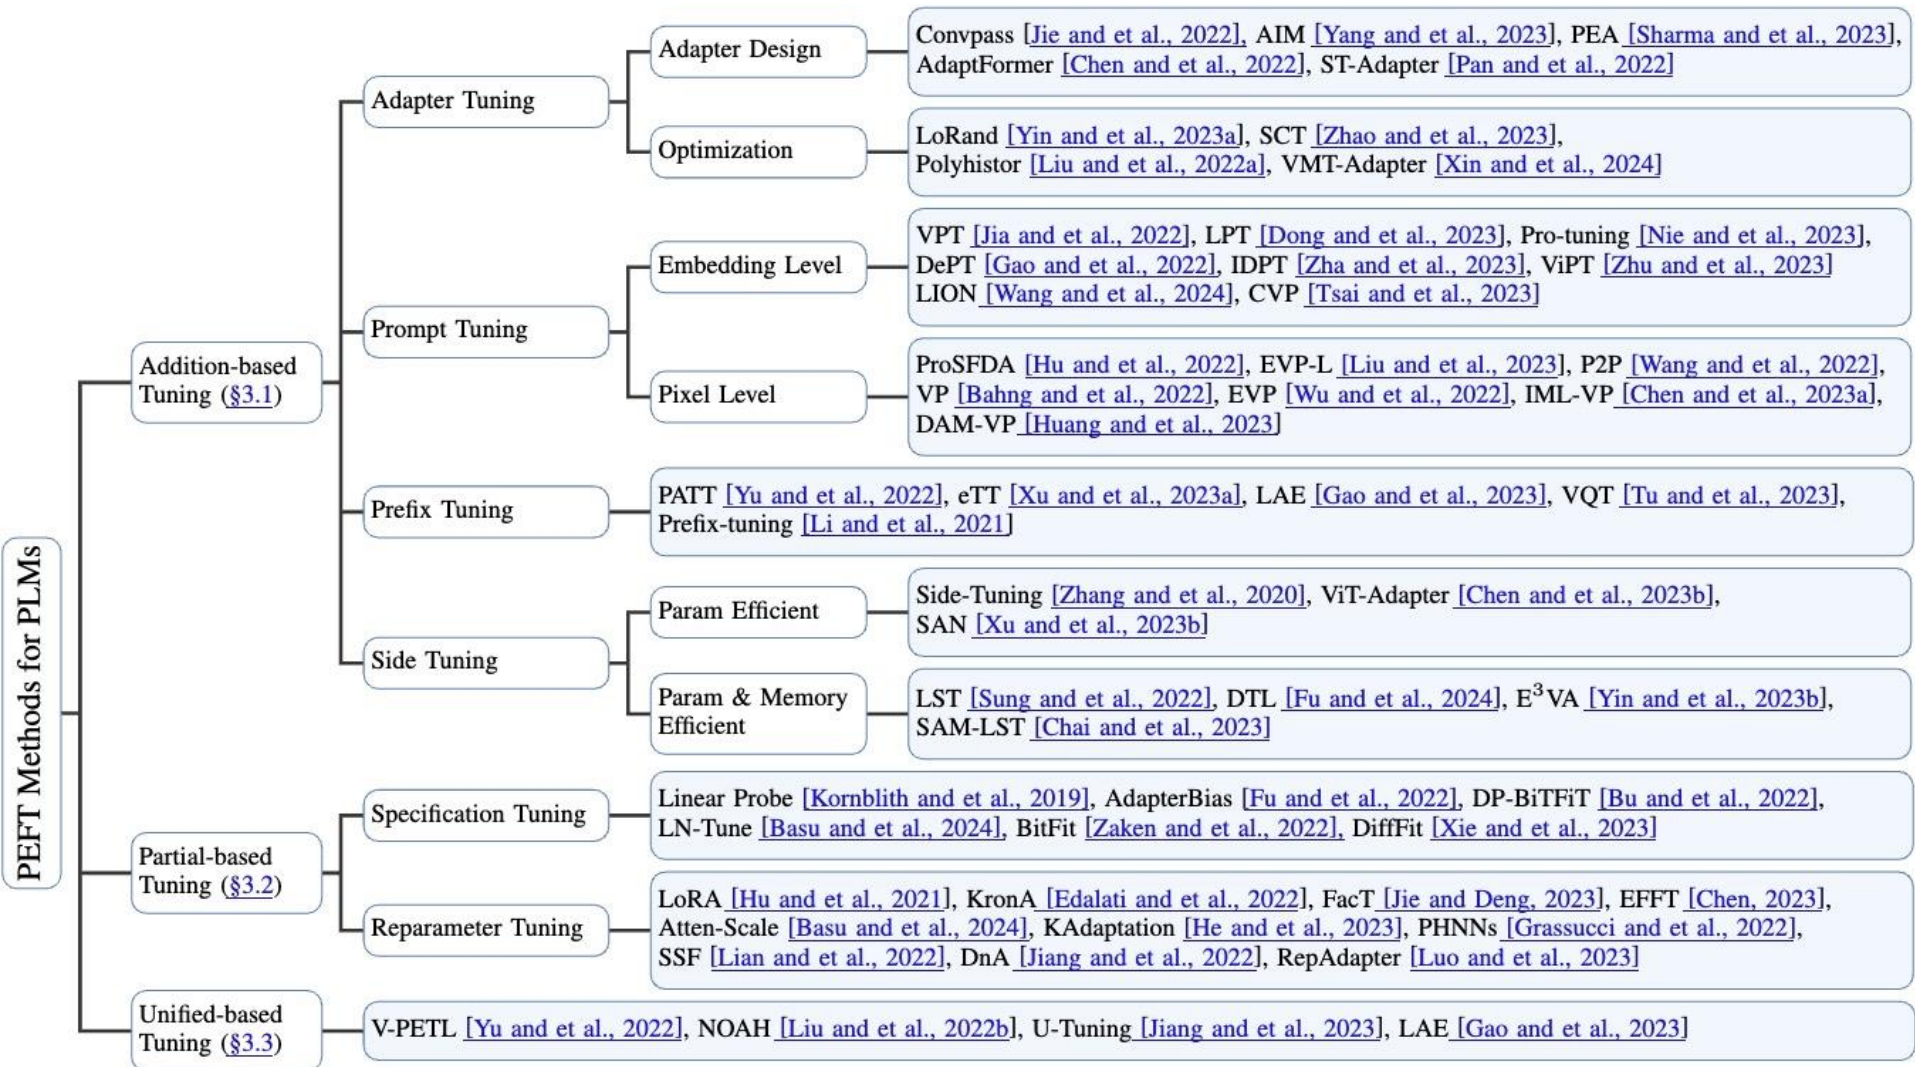


Figure 1: Taxonomy of Parameter-Efficient Fine-Tuning Methods for Pre-trained Vision Models.

Xin, Yi, et al. "Parameter-efficient fine-tuning for pre-trained vision models: A survey." *arXiv preprint arXiv:2402.02242* (2024).

# Tools

---

- <https://github.com/huggingface/peft>



## State-of-the-art Parameter-Efficient Fine-Tuning (PEFT) methods

Fine-tuning large pretrained models is often prohibitively costly due to their scale. Parameter-Efficient Fine-Tuning (PEFT) methods enable efficient adaptation of large pretrained models to various downstream applications by only fine-tuning a small number of (extra) model parameters instead of all the model's parameters. This significantly decreases the computational and storage costs. Recent state-of-the-art PEFT techniques achieve performance comparable to fully fine-tuned models.

PEFT is integrated with Transformers for easy model training and inference, Diffusers for conveniently managing different adapters, and Accelerate for distributed training and inference for really big models.

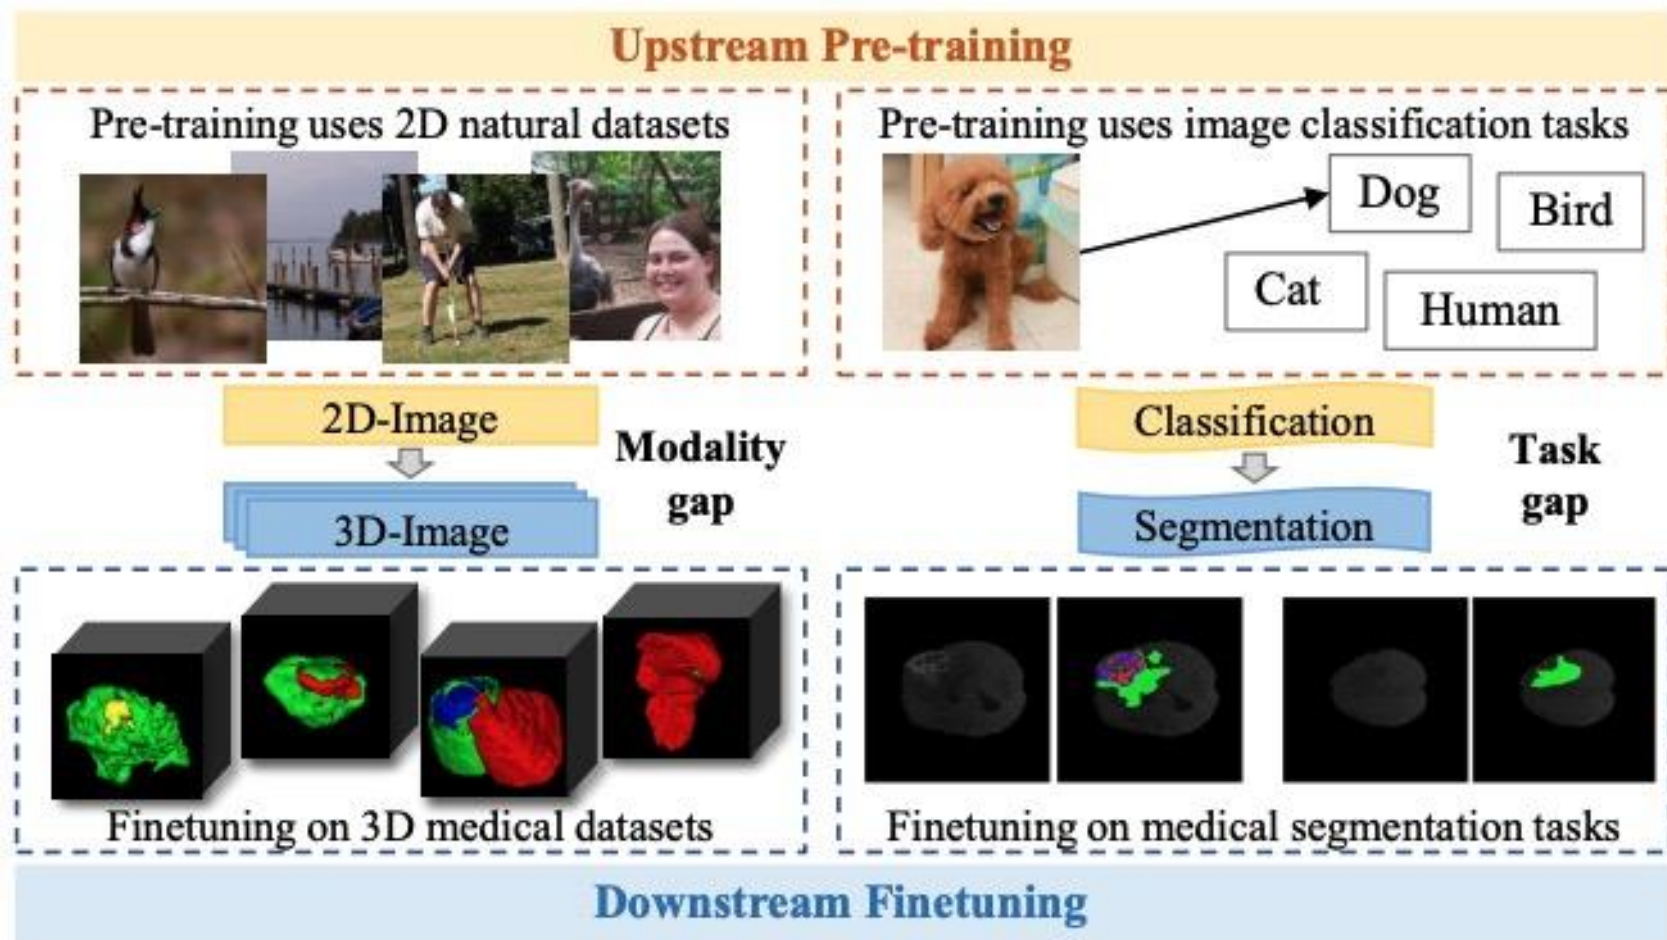
---

# Parameter Efficient Fine Tuning for Medical Volumetric Segmentation




# Parameter Efficient Fine Tuning for Medical Volumetric Segmentation

- Challenges



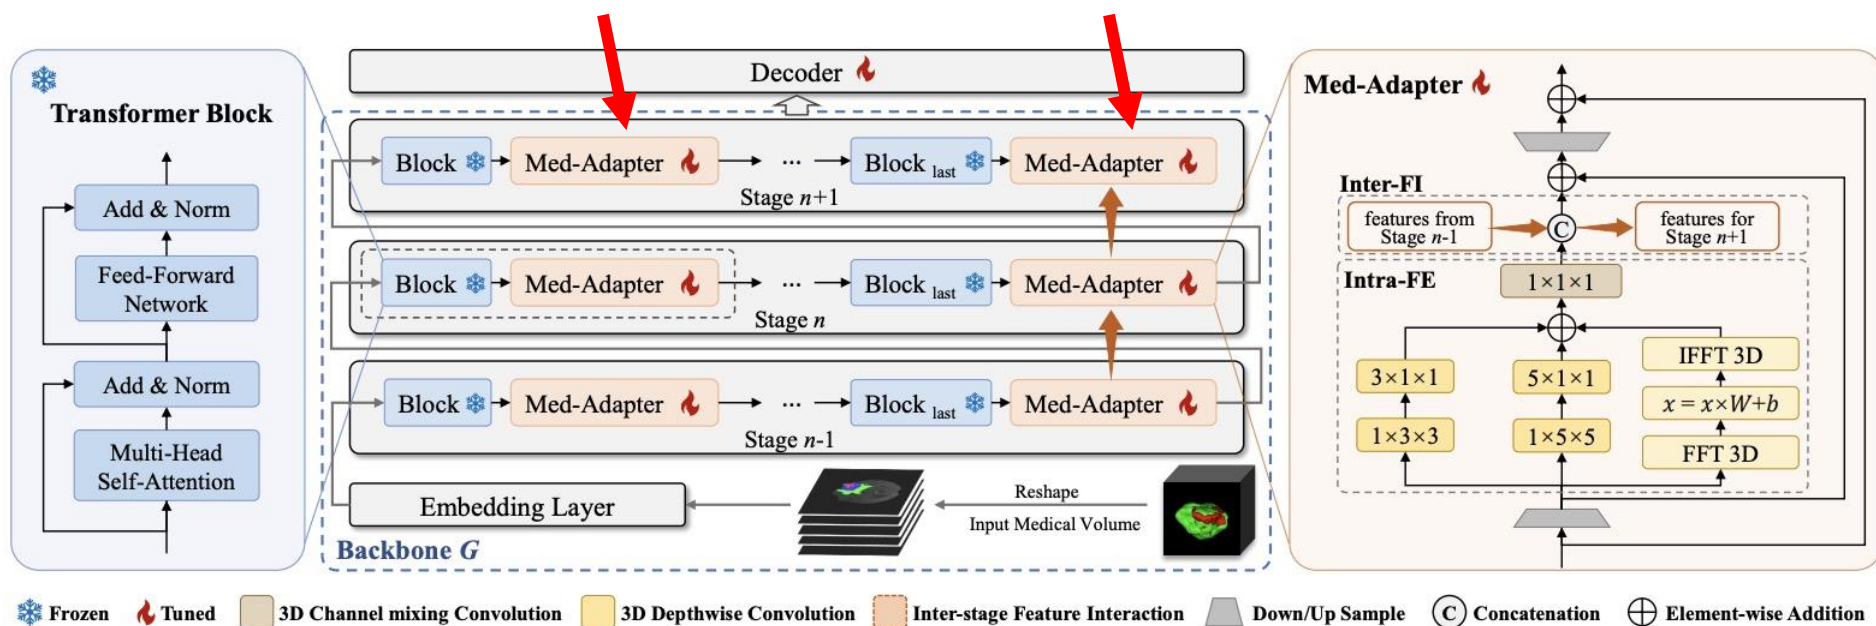
# Parameter Efficient Fine Tuning for Medical Volumetric Segmentation

- Goal
  - Empower a 2D Transformer model pre-trained on natural images to gain the capability of **spatial and temporal feature modeling** for medical volumetric data in a **parameter-efficient manner**

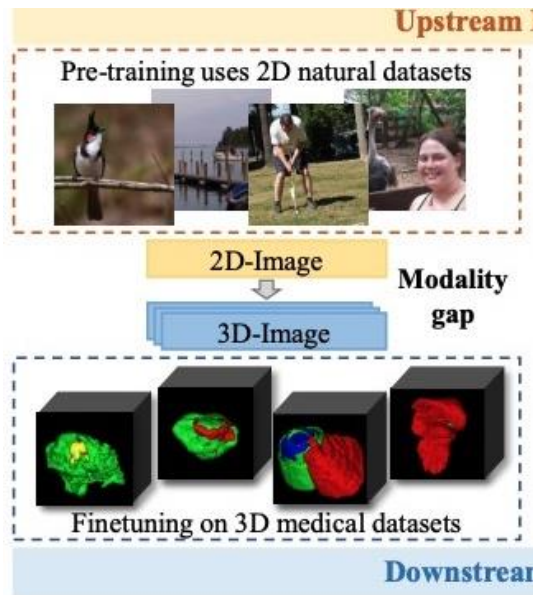


# Proposed Method – Med-Tuning

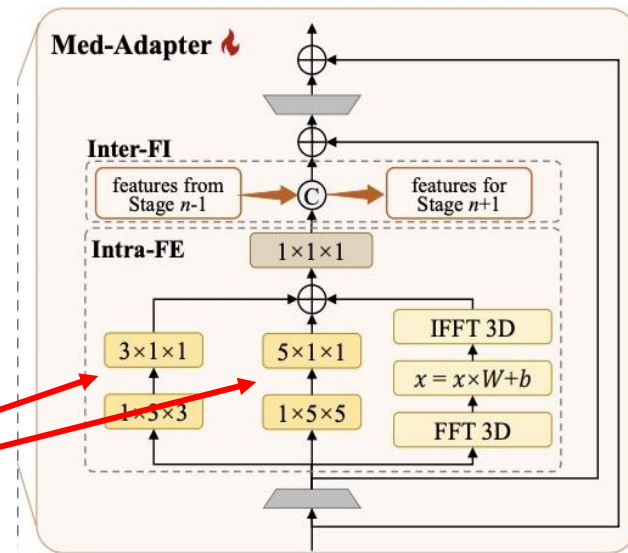
- Follow the overall structure of the adapter module with tailored designs



# Proposed Method – Med-Tuning



Capture the correlations among neighboring slices

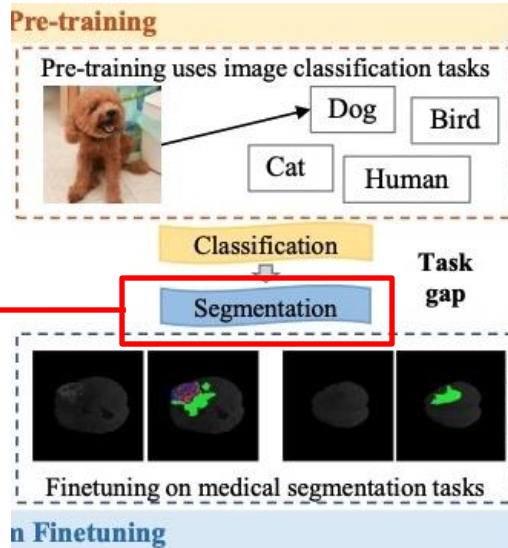


Frozen 
 Tuned 
 3D Channel mixing Convolution 
 3D Depthwise Convolution 
 Inter-stage Feature Interaction 
 Down/Up Sample 
 C Concatenation 
  $\oplus$  Element-wise Addition

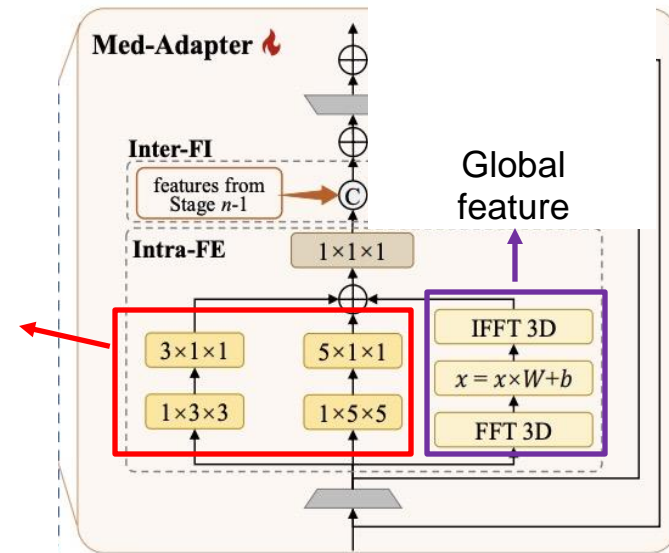
# Proposed Method – Med-Tuning

Dense prediction task (assign a label for each pixel)

Both local (fine-grained) feature representations and global semantic information are important



Multi-scale feature representation



Frozen 
 Tuned 
 3D Channel mixing Convolution 
 3D Depthwise Convolution 
 Inter-stage Feature Interaction 
 Down/Up Sample 
 Concatenation 
 Element-wise Addition



# Results - BraTS

---

MRI brain tumor  
segmentation

ViT + UPerNet	Tuned Params (M)	Inserted Params (M)	BraTS2019						BraTS2020					
			Dice (%) ↑			Hausdorff (mm) ↓			Dice (%) ↑			Hausdorff (mm) ↓		
			ET	WT	TC	ET	WT	TC	ET	WT	TC	ET	WT	TC
Scratch	100.849	-	64.96	83.03	71.34	7.635	10.602	10.942	65.80	83.72	72.01	32.475	10.060	21.467

# Results

---

## MRI brain tumor segmentation

ViT + UPerNet	Tuned Params (M)	Inserted Params (M)	BraTS2019						BraTS2020					
			Dice (%) ↑			Hausdorff (mm) ↓			Dice (%) ↑			Hausdorff (mm) ↓		
			ET	WT	TC	ET	WT	TC	ET	WT	TC	ET	WT	TC
Scratch	100.849	-	64.96	83.03	71.34	7.635	10.602	10.942	65.80	83.72	72.01	32.475	10.060	21.467
Full	100.849	-	68.49	85.56	75.12	6.672	7.878	10.525	69.12	85.90	75.29	34.428	7.315	17.093



# Results

## MRI brain tumor segmentation

ViT + UPerNet	Tuned Params (M)	Inserted Params (M)	BraTS2019						BraTS2020					
			Dice (%) ↑			Hausdorff (mm) ↓			Dice (%) ↑			Hausdorff (mm) ↓		
			ET	WT	TC	ET	WT	TC	ET	WT	TC	ET	WT	TC
Scratch	100.849	-	64.96	83.03	71.34	7.635	10.602	10.942	65.80	83.72	72.01	32.475	10.060	21.467
Full	100.849	-	68.49	85.56	75.12	6.672	7.878	10.525	69.12	85.90	75.29	34.428	7.315	17.093
Head	15.007	-	65.71	84.19	74.77	6.128	7.505	7.864	66.03	84.50	74.47	37.805	7.474	14.150
VPT-Shallow [26]	15.015	0.008	66.02	84.72	75.84	6.114	7.506	8.471	66.52	84.82	75.46	37.765	7.465	13.531
VPT-Deep [26]	15.100	0.092	67.01	85.14	76.80	6.064	7.717	7.648	67.69	85.28	76.59	31.772	7.737	<b>10.621</b>

# Results

MRI brain tumor  
segmentation

ViT + UPerNet	Tuned Params (M)	Inserted Params (M)	BraTS2019						BraTS2020					
			Dice (%) ↑			Hausdorff (mm) ↓			Dice (%) ↑			Hausdorff (mm) ↓		
			ET	WT	TC	ET	WT	TC	ET	WT	TC	ET	WT	TC
Scratch	100.849	-	64.96	83.03	71.34	7.635	10.602	10.942	65.80	83.72	72.01	32.475	10.060	21.467
Full	100.849	-	68.49	85.56	75.12	6.672	7.878	10.525	69.12	85.90	75.29	34.428	7.315	17.093
Head	15.007	-	65.71	84.19	74.77	6.128	7.505	7.864	66.03	84.50	74.47	37.805	7.474	14.150
VPT-Shallow [26]	15.015	0.008	66.02	84.72	75.84	6.114	7.506	8.471	66.52	84.82	75.46	37.765	7.465	13.531
VPT-Deep [26]	15.100	0.092	67.01	85.14	76.80	6.064	7.717	7.648	67.69	85.28	76.59	31.772	7.737	<b>10.621</b>
Adapter [21]	18.567	3.560	68.30	85.37	77.05	<b>5.501</b>	7.636	7.986	68.58	85.77	77.00	32.626	8.172	16.183
AdaptFormer [11]	16.197	1.190	65.88	84.34	74.77	6.652	8.204	8.430	65.52	84.14	74.28	41.026	8.393	14.778
Pro-tuning [35]	19.812	4.805	67.18	85.32	76.51	5.805	7.073	7.564	67.28	85.57	76.58	40.434	7.000	12.865
ST-Adapter [37]	22.118	7.110	69.18	86.27	79.18	6.077	6.939	<b>6.778</b>	68.60	86.55	<b>79.52</b>	34.060	6.790	12.770



# Results

MRI brain tumor  
segmentation

ViT + UPerNet	Tuned Params (M)	Inserted Params (M)	BraTS2019						BraTS2020					
			Dice (%) ↑			Hausdorff (mm) ↓			Dice (%) ↑			Hausdorff (mm) ↓		
			ET	WT	TC	ET	WT	TC	ET	WT	TC	ET	WT	TC
Scratch	100.849	-	64.96	83.03	71.34	7.635	10.602	10.942	65.80	83.72	72.01	32.475	10.060	21.467
Full	100.849	-	68.49	85.56	75.12	6.672	7.878	10.525	69.12	85.90	75.29	34.428	7.315	17.093
Head	15.007	-	65.71	84.19	74.77	6.128	7.505	7.864	66.03	84.50	74.47	37.805	7.474	14.150
VPT-Shallow [26]	15.015	0.008	66.02	84.72	75.84	6.114	7.506	8.471	66.52	84.82	75.46	37.765	7.465	13.531
VPT-Deep [26]	15.100	0.092	67.01	85.14	76.80	6.064	7.717	7.648	67.69	85.28	76.59	31.772	7.737	<b>10.621</b>
Adapter [21]	18.567	3.560	68.30	85.37	77.05	<b>5.501</b>	7.636	7.986	68.58	85.77	77.00	32.626	8.172	16.183
AdaptFormer [11]	16.197	1.190	65.88	84.34	74.77	6.652	8.204	8.430	65.52	84.14	74.28	41.026	8.393	14.778
Pro-tuning [35]	19.812	4.805	67.18	85.32	76.51	5.805	7.073	7.564	67.28	85.57	76.58	40.434	7.000	12.865
ST-Adapter [37]	22.118	7.110	69.18	86.27	79.18	6.077	6.939	<b>6.778</b>	68.60	86.55	<b>79.52</b>	34.060	6.790	12.770
Ours	17.853 (17.70%)	2.846 (2.82%)	<b>70.53</b> (+2.04)	<b>86.58</b> (+1.02)	<b>79.35</b> (+4.23)	5.862 (-0.810)	<b>6.224</b> (-1.654)	6.947 (-3.578)	<b>70.69</b> (+1.57)	<b>86.69</b> (+0.79)	79.36 (+4.07)	<b>28.643</b> (-5.785)	<b>6.198</b> (-1.117)	15.045 (-2.048)



# Results

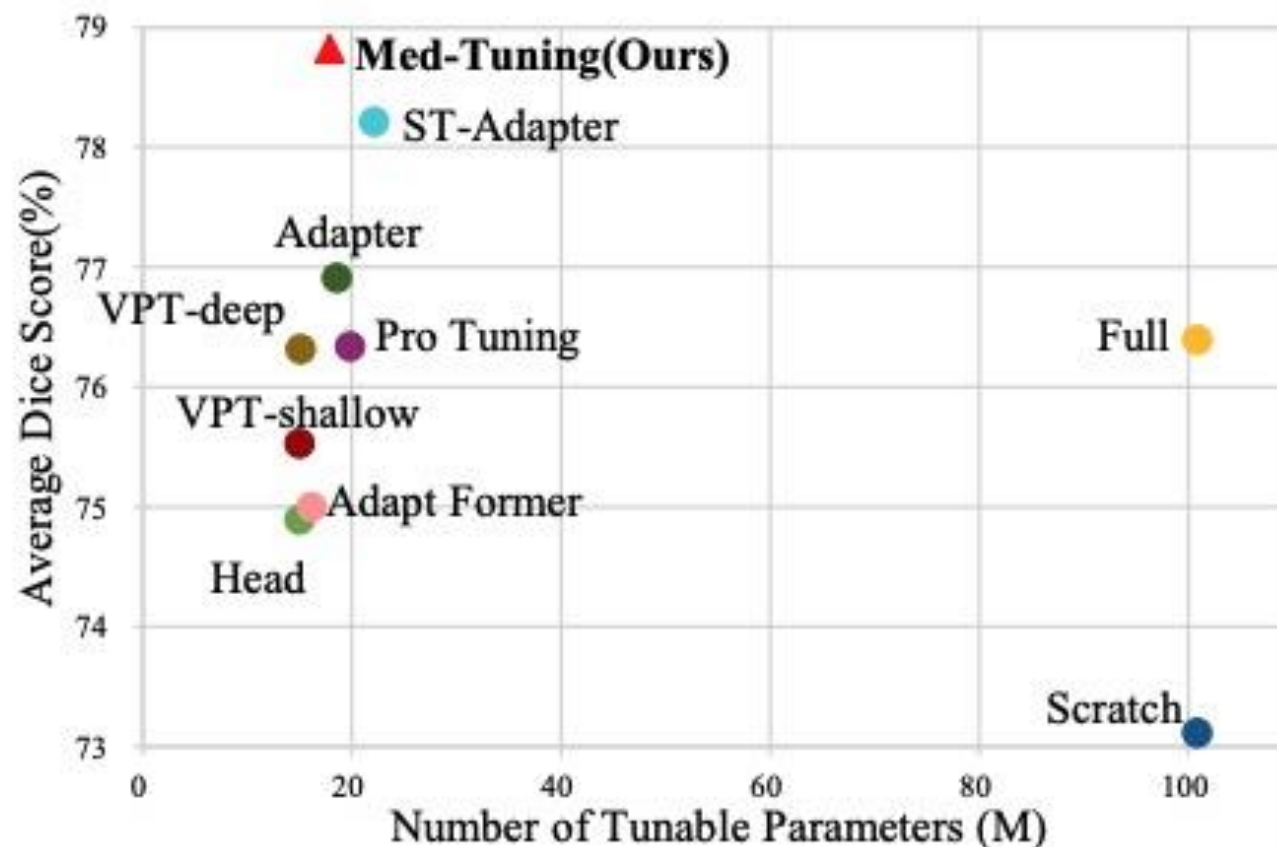


Figure 2. Comparison with previous PETL methods in terms of the trade-off between the number of tuned parameters and segmentation accuracy. The backbone ViT-B/16 is pre-trained on ImageNet-21k and fine-tuned on BraTS2019 dataset.

# Results - KiTS

Kidney tumor segmentation on CT

Swin-UNet	Tuned Params(M)	Inserted Params(M)	Dice (%) $\uparrow$			ViT + UPerNet	Tuned Params(M)	Inserted Params(M)	Dice (%) $\uparrow$		
			Kidney	Tumor	Composite				Kidney	Tumor	Composite
Scratch	27.154	-	94.33	61.10	77.71	Scratch	100.849	-	88.01	46.53	67.27
Full	27.154	-	94.68	62.13	78.40	Full	100.849	-	87.32	47.34	67.33
Head	6.752	-	91.95	53.93	72.94	Head	15.007	-	87.35	42.85	65.10
VPT-Shallow [26]	6.753	0.001	91.72	54.86	73.29	VPT-Shallow	15.015	0.008	86.91	41.67	64.29
VPT-Deep [26]	6.780	0.029	91.53	53.41	72.47	VPT-Deep	15.100	0.092	88.01	46.45	67.23
Adapter [21]	7.541	0.790	93.02	57.15	75.08	Adapter	18.567	3.560	89.75	49.03	69.39
AdaptFormer [11]	7.124	0.372	93.74	59.79	76.77	AdaptFormer	16.197	1.190	87.62	44.46	66.04
Pro-tuning [35]	8.359	1.607	90.34	51.19	70.77	Pro-tuning	19.812	4.805	89.44	48.32	68.88
ST-Adapter [37]	8.328	1.577	92.97	57.33	75.15	ST-Adapter	22.118	7.110	90.33	61.29	75.81
<b>Ours</b>	<b>7.489</b> <b>(27.58%)</b>	<b>0.738</b> <b>(2.72%)</b>	<b>95.69</b> <b>(+1.01)</b>	<b>70.14</b> <b>(+8.02)</b>	<b>82.92</b> <b>(+4.52)</b>	<b>Ours</b>	<b>17.853</b> <b>(17.70%)</b>	<b>2.846</b> <b>(2.82%)</b>	<b>91.52</b> <b>(+4.20)</b>	<b>64.47</b> <b>(+17.13)</b>	<b>78.00</b> <b>(+10.67)</b>



# Visual Results

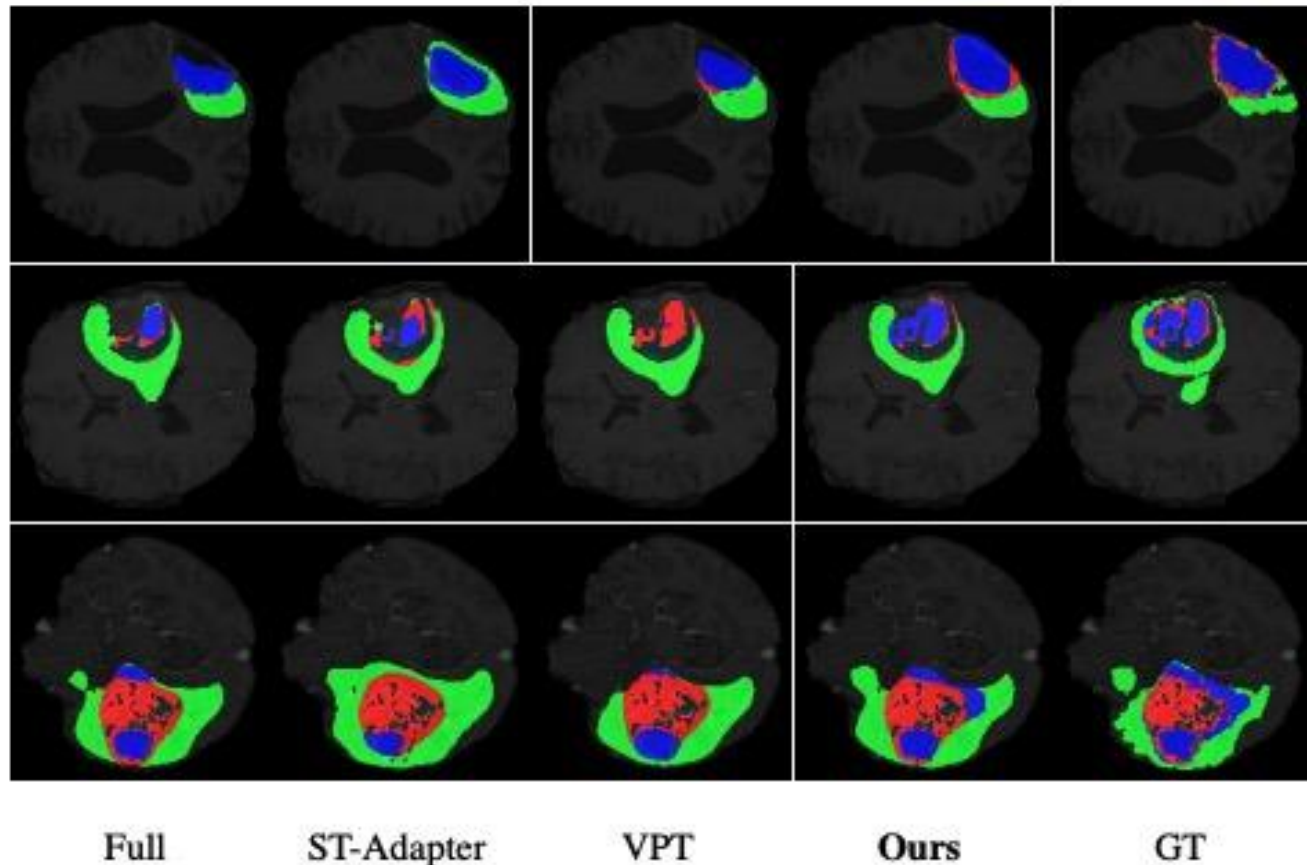


Figure 4. The visual comparison of segmentation results on BraTS 2019 dataset. The blue, red and green regions denote the enhancing tumors, non-enhancing tumors, and peritumoral edema. Full and GT denote full fine-tuning and ground truth.



# Visual Results




Figure 5. The visual comparison of segmentation results on KiTS 2019 dataset. The **red** and **green** regions denote the kidneys and kidney tumors. Full, GT denote full finetuning and ground truth.



---

Thank you!

Question?